


Learning Probabilistic Automata



Mathias M. Andersen

Kent M. Caspersen

Anders R. Nielsen

Juraj Kolčák

Guillaume Grippari

Theis Christensen

MODELS

n grams \approx Markov chains \leq DPFA... \leq

\leq HMM \approx PFA \leq

\leq Multiplicity automata

PAutomaC

- ▶ 2012 online PFA learning competition.
- ▶ All results and training data published after competition finished.
- ▶ Large amount of data available artificially generated by several types of models (Markov chains, DPFA, PFA, HMM).
- ▶ The generator models vary in number of states/symbols and density of the transitions and emissions providing for a large scale of different data.

PROBLEM DEFINITION

“Based on the *PAutomaC* competition data, how can we learn not only the probabilistic parameters but the the structure of an HMM as well.”

HIDDEN MARKOV MODELS

DEFINITION

Alphabet of observable symbols $\Sigma = \{\sigma_1, \dots, \sigma_m\}$.

Set of hidden states $S = \{s_1, \dots, s_n\}$.

Sequence of observable symbols $\mathbf{O} = (o_1, \dots, o_T) = \Sigma^T$.

Sequence of hidden states $\mathbf{Q} = (q_1, \dots, q_T) = S^T$.

$$\lambda = (\mathbf{A}, \mathbf{B}, \boldsymbol{\pi})$$

Transition matrix \mathbf{A} : $\forall s, r \in S : a_{sr} = P(q_{t+1} = r | q_t = s)$

Emission matrix \mathbf{B} : $\forall s \in S, \forall \sigma \in \Sigma : b_s(\sigma) = P(o_t = \sigma | q_t = s)$

Initial probability distribution $\boldsymbol{\pi}$: $\forall s \in S : \pi_s = P(q_1 = s)$

EVALUATION

Evaluation signal $\mathbf{O} = (o_1, \dots, o_T)$ and model λ .

$$\begin{aligned} P(\mathbf{O}|\lambda) &= \sum_{\mathbf{Q} \in \mathcal{Q}_{\lambda}^T} P(\mathbf{O}|\mathbf{Q}, \lambda) P(\mathbf{Q}|\lambda) \\ &= \sum_{\mathbf{Q} \in \mathcal{Q}_{\lambda}^T} \pi_{q_1} b_{q_1}(o_1) a_{q_1 q_2} b_{q_2}(o_2) \dots a_{q_{T-1} q_T} b_{q_T}(o_T) \end{aligned}$$

Where \mathcal{Q}_{λ}^T is a set of all walks through the hidden state space of λ of length T .

FORWARD-BACKWARD PROCEDURE

Forward variable:

$$\forall s \in S : \alpha_1(s) = \pi_s b_s(o_1)$$

$$\forall s \in S, t \in \{2, \dots, T\} : \alpha_t(s) = \sum_{r \in S} (\alpha_{t-1}(r) a_{rs}) b_s(o_t)$$

Backward variable:

$$\forall s \in S : \beta_T(s) = 1$$

$$\forall s \in S, t \in \{1, \dots, T-1\} : \beta_t(s) = \sum_{r \in S} (a_{sr} b_r(o_{t+1}) \beta_{t+1}(r))$$

$$\sum_{s \in S} \alpha_T(s) = P(\mathbf{O}|\lambda) = \sum_{s \in S} \beta_1(s)$$

BAUM-WELCH ALGORITHM

$$\begin{aligned}\forall t \in \{1, \dots, T-1\}, \forall s, r \in S : \xi_t(s, r) &= P(q_t = s, q_{t+1} = r | \mathbf{O}, \lambda) = \\ &= \frac{\alpha_t(i) a_{sr} b_r(o_{t+1}) \beta_{t+1}(r)}{\sum_{u \in S} \sum_{v \in S} (\alpha_t(u) a_{uv} b_v(o_{t+1}) \beta_{t+1}(v))}\end{aligned}$$

$$\forall t \in \{1, \dots, T\}, \forall s \in S : \gamma_t(s) = P(q_t = s | \mathbf{O}, \lambda) = \sum_{r \in S} \xi_t(s, r)$$

BAUM-WELCH ALGORITHM

$$\forall s \in S : \overline{\pi_s} = \gamma_1(s)$$

$$\forall s, r \in S : \overline{a_{sr}} = \frac{\sum_{t=1}^{T-1} \xi_t(s, r)}{\sum_{t=1}^{T-1} \sum_{u \in S} \xi_t(s, u)}$$

$$\forall s \in S, \forall \sigma \in \Sigma : \overline{b_s(\sigma)} = \frac{\sum_{t \in \mathcal{T}_{\mathbf{O}}(\sigma)} \gamma_t(s)}{\sum_{t=1}^T \gamma_t(s)}$$

$$\mathcal{T}_{\mathbf{O}}(\sigma) = \{t \in \{1, \dots, T\} | o_t = \sigma\}$$

$$P(\mathbf{O} | \overline{\lambda} = (\overline{\mathbf{A}}, \overline{\mathbf{B}}, \overline{\pi})) \geq P(\mathbf{O} | \lambda = (\mathbf{A}, \mathbf{B}, \pi))$$

HMM vs PFA

- ▶ HMMs and PFAs are mutually convertible between each other.
- ▶ PFAs often utilise stopping probabilities (also used for *PAutomaC* models).
- ▶ Distribution defined by PFA with stopping probabilities: $P(\Sigma^*)$ against the HMM: $\forall n \in \mathbb{N} : P(\Sigma^n)$.
- ▶ A possible solution: introduce a new “stopping” symbol $x \notin \Sigma$ and create an HMM over the alphabet $\bar{\Sigma} = \Sigma \cup \{x\}$. End the signal once the new symbol x is reached.

HMM DENSITY

- ▶ Empirical results show that real world entities display “sparse” behaviour.
- ▶ Current state-of-the-art methods (Baum-Welch for HMMs) require the user to fully specify the amount of states and structure of the transition graph.
- ▶ Sparse transition matrix can also provide for a computational speedup.

OVERFITTING

Example:

We want to learn the model that generated some training data:

Training data

2135
4313
1325
2213
4133

Learned model

2135
4313
2135
2213
4133

Learned model

2, 2135
1, 4313
1, 2213
1, 4133

LL of training data: $\sum_O \log P(O \mid \lambda)$

What about other data generated by the model?

New sequences	Learned model
4135	2, 2135
1322	1, 4313
4213	1, 2213
5135	1, 4133
1345	

LL of new data?

► 0?

When does overfitting happen?


- ▶ Learning a too complex model
 - ▶ Models all the noise
- ▶ Training data too little
 - ▶ May not reflect model good enough

Assume a very general pattern exists:

Training data

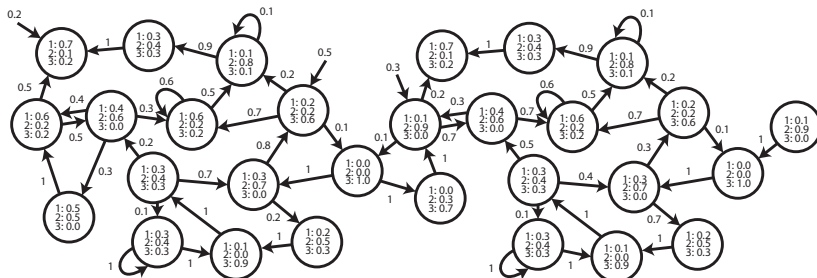
2 <u>1</u> 35
43 <u>1</u> 3
<u>1</u> 325
22 <u>1</u> 3
4 <u>1</u> 33

1 always followed by 3:



When learning a simple model:


- ▶ Cannot express all noise
- ▶ May benefit more from describing general patterns



When learning a complex model:

- May be able to describe all noise
 - So why not do it?

How do we compare models / algorithms?



AVOIDING UNDERFLOW

$$P(X) = p_1 p_2 p_3 p_4 p_5 p_6 p_7 p_8 p_9 \cdots$$

$$0 < p_x < 1$$

Underflow when $P(X) = 0$

Visual Studio, C#: Double precision floating point

- ▶ $0.01^{200} = 0$
- ▶ $0.02^{200} = 0$

Probability of a particular sequence can be very small:

- ▶ 23 symbols, uniformly distributed: $\frac{1}{23}^{238} = 0$
- ▶ What about rare symbols?
- ▶ What about probability of 1000 sequences?

Solution: Convert to log-space

$$\log (a \cdot b) = \log a + \log b$$

Since $a < b \iff \log a < \log b$:

We can often stay in Log-space

$$\log(0.01^{200}) = \log 0.01 + \log 0.01 + \cdots = -1381.55$$

$$\log(0.02^{200}) = \log 0.02 + \log 0.02 + \cdots = -1173.60$$

ALGORITHMS

STATIC VS. DYNAMIC




Figure: The algorithms used in our experiments

SPARSE BAUM-WELCH

- ▶ Creates HMM with n states and m symbols
- ▶ All parameters are initialized randomly
- ▶ Constraint: Each state has exactly $\log(n)$ transitions
- ▶ Other transitions set to zero
- ▶ Trained using Baum-Welch until convergence

GREEDY EXTEND: SETUP

- ▶ Works by adding states to the graph in iterations
- ▶ Starts as a single node with initial probability 1, random emission probabilities and a single transition to itself.

GREEDY EXTEND: ITERATIONS

1. Repeat until convergence
2. $G' = (V(G) \cup \{y'\}, E(G))$
3. Randomly choose a set X of $\log|V(G')|$ nodes from $V(G')$
4. $\forall x \in X$ add transitions (x, y') and (y', x) to $E(G')$ with random probabilities
5. Normalize G'
6. if $LL(BW^\beta(G')) > LL(G)$, let $G = BW^\beta(G')$

STATE SPLITTING: OVERALL APPROACH

1. Identify a set of states \mathcal{W} to split
2. Split all states in \mathcal{W} using a mechanic
3. Run Baum-Welch for β iterations

STATE SPLITTING: SPLITTING MECHANICS

- ▶ Clone Split
 - ▶ Makes a copy of the chosen state
 - ▶ Problem: BW unable to distinguish between clone and original
 - ▶ Alternative: Randomize clones probabilities.
- ▶ Distribution Split
 - ▶ Only splits if Transition or Emission probabilities are uniform.
 - ▶ Copies the emission probabilities from the original to the new state
 - ▶ Randomizes transition probabilities on the new state
 - ▶ Problem: Algorithm can get stuck (splits after 10 iterations)

STATE SPLITTING: IDENTIFICATION HEURISTICS

- ▶ The Heuristics compute a score ς that is used to choose which states to split
- ▶ Gamma Heuristic
 - ▶ Assign ς based on the number of times the state is visited when generating the sequence
 - ▶ $\forall i \in \{1, \dots, n\} : \varsigma S_i = \sum_{O \in D} \sum_{t=1}^T \gamma_t(S_i)$

STATE SPLITTING: IDENTIFICATION HEURISTICS

► Viterbi Heuristic

1. Compute $Q = \mathcal{V}_G(O)$ foreach signal $O \in D$
2. Foreach state $s \in S$ determine its significant positions in Q
3. $\forall s \in S, \forall O \in D$ compute $\hat{\varsigma}_O(s) = \frac{\sum_{t \in \tau_{s,\lambda}^O} b_s(o_t)}{|\tau_{s,\lambda}^O|}$
4. Compute $\forall s \in S: \varsigma(s) = \sum_{O \in D} P(Q|O, \lambda) \hat{\varsigma}_O(s)$

EDGE CUTTING & STATE REMOVAL

- ▶ Edge cutting
 - ▶ Strict Edge cutting
 - ▶ Threshold Edge cutting
- ▶ State Removal

CHOSEN ALGORITHM

The algorithm we chose for the experiments had the following characteristics

- ▶ Splitting Mechanic: Distribution Split
- ▶ Identification: Gamma Heuristic
- ▶ β value: 10

TEST ENVIRONMENT

TEST ENVIRONMENT

- ▶ Benchmark
- ▶ Data
- ▶ PautomacEvaluator
- ▶ Learner
- ▶ Models

SELECTING DATASET

SELECTING DATASET

- ▶ Sparsity
- ▶ Transition Density

EXPERIMENT PARAMETER

EXPERIMENT PARAMETER

- ▶ Static Learners
- ▶ Training Sequences
- ▶ Baum-Welch Threshold
- ▶ Greedy Extend


EXPERIMENTS

EXPERIMENTS

- ▶ Training - 5000 sequences
- ▶ Validation - 5000 sequences
- ▶ States - 10 to 100
- ▶ How does problem characteristics affect prediction accuracy
- ▶ Speed comparison


STATE SPACE

Data set: 6, States: 19




STATE SPACE

Data set: 23, States: 33




STATE SPACE

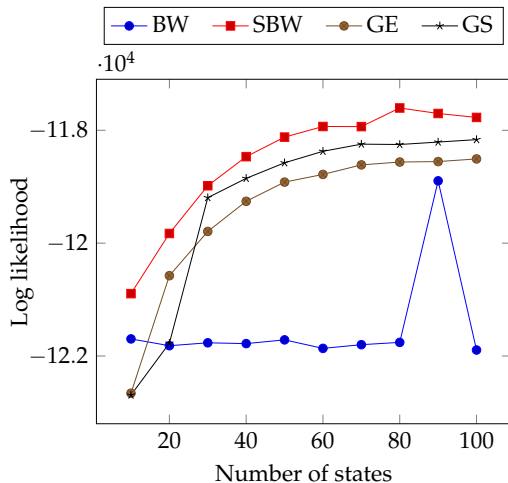
Data set: 41, States: 54



STATE SPACE


Data set: 1, States: 64

TRANSITION SPARSITY

Dataset: 36, Sparsity: 7.4%


TRANSITION SPARSITY

Dataset: 8, Sparsity: 16.8%




TRANSITION SPARSITY

Dataset: 43, Sparsity: 40.2%




TRANSITION SPARSITY


Dataset: 37, Sparsity: 50%




ALPHABET SIZE

Dataset: 32, Symbols: 4

ALPHABET SIZE


Dataset: 8, Symbols: 8

ALPHABET SIZE

Dataset: 10, Symbols: 11

ALPHABET SIZE

Dataset: 35, Symbols: 20



SPEED COMPARISON




Figure: Results achieved in a time scope of eight hours.

SPEED COMPARISON




Figure: Running time comparison between Baum-Welch and Sparse Baum-Welch Learners

RESULTS

- ▶ 3 data sets: 6, 23, 35
- ▶ Parameters chosen based on the experiments presented earlier
- ▶ 5000 sequences
- ▶ 10 initial states, step size 10 for static algorithms
- ▶ 10 initial states, step size 5 for dynamic algorithms

RESULTS

Data set	GE	BW	SBW	GS	Goal
6	115.53	122.60	129.64	114.77	66.98
23	26.30	26.16	26.08	26.19	18.41
35	49.48	43.36	50.79	44.44	33.78

Table: The best scores of each algorithm on the three data sets.

RESULTS

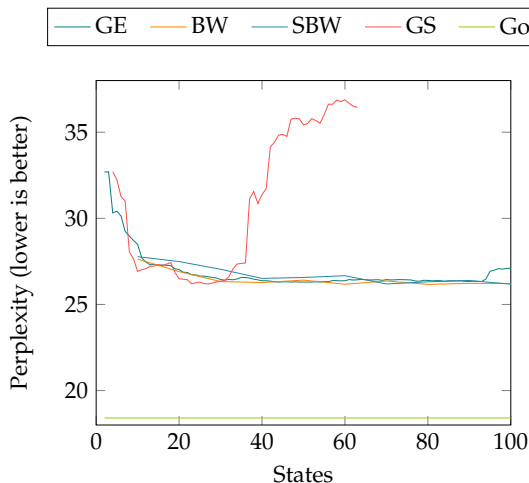


Figure: The different algorithm's score on *PAutomaC*'s 23rd data set, according to the perplexity measure used in the competition.

CONCLUSION

CONCLUSION

- ▶ Analysis of different approaches using a sparse transition matrix in the BW algorithm
- ▶ Dynamic algorithms: about as good results as BW when not better
- ▶ Sparse matrix ensured by adding states iteratively

CONCLUSION: PROS AND CONS

Algorithm	Advantages	Drawbacks
GE	<ul style="list-style-type: none">- Rather fast- Good results	<ul style="list-style-type: none">- Usually needs more states to start to compare
GS	<ul style="list-style-type: none">- Splitting heuristic- Best results, but...	<ul style="list-style-type: none">- Runs on a dense matrix- ...unusual behaviour
SBW	<ul style="list-style-type: none">- Computational speed-up- Room for improvement	<ul style="list-style-type: none">- Results slightly worse than BW- No speed-up on some data sets

CONCLUSION: EXPERIMENTS CONDUCTED

- ▶ 3 types of experiment
- ▶ 4 algorithms

CONCLUSION: FIRST EXPERIMENTS

- ▶ Extremely varied behaviour on different data sets
- ▶ Three available parameters for results analysis
- ▶ Hypotheses, but no definite clear pattern
- ▶ In general, results on par with BW (sometimes outperforming it)

CONCLUSION: SECOND EXPERIMENTS

- ▶ Speed comparison between: BW, SBW, GE
- ▶ 8-hour run
- ▶ GE dominating both in speed and in score
- ▶ SBW unstable: random matrix causing non-deterministic behaviour (suggests building a data-derived one instead)

CONCLUSION: THIRD EXPERIMENTS

- ▶ Comparison using PAutomaC scores
- ▶ Unsatisfactory results, but...
- ▶ ...due to a different paradigm used: no stopping probabilities have been used

CONCLUSION: FUTURE WORK

- ▶ Eliminating possible noise
- ▶ Explain the multitude of behaviours, validating or disproving our hypotheses
- ▶ Running on more states
- ▶ Comparing with the PAutomaC scores, using the same paradigm
- ▶ Optimising the sparse transition matrices used in the algorithms
- ▶ Making sure underflow and overfitting are avoided
- ▶ Looking more in-depth into the other methods we have put aside in this project